

SETHURAM REDDY LANKALA

Houston, TX (Open to Relocation) • sethuramreddy.lankala@gmail.com • +1 (202) 213-2127 • LinkedIn • GitHub • Portfolio

SUMMARY

Founding AI / Backend Engineer with 3+ years across startup product ownership, enterprise systems, and applied AI. Built a recruiting platform from customer discovery through production and a deployed seven-agent research application. Strong in Python services, data ingestion, retrieval, agent orchestration, evaluation, and turning ambiguous workflows into working products. Open to relocating to San Francisco.

EXPERIENCE

Data Science Fellow

Sep 2026 - Present

ChiEAC Data Science Alliance

Remote

- Translate stakeholder needs into scoped Data, AI, and Agentic AI workflows, reproducible Python implementations, evaluation plans, and technical documentation.
- Assess data quality, model behavior, failure cases, and deployment constraints; prioritize experiments and communicate trade-offs to stakeholders.

Graduate AI Researcher - RAG Systems & Evaluation

Jan 2026 - Jun 2026

The George Washington University (Advisor: Prof. Nan Wu)

Washington, DC

- Transformed 500+ heterogeneous C/C++/TCL sources into validated retrieval datasets; engineered ingestion, metadata, chunking, embedding, semantic search, and grounded-generation pipelines.
- Built reproducible benchmarks and failure taxonomies to compare retrieval, prompt, and model changes on accuracy, relevance, latency, reliability, and cost.

Chief Technology Officer & Founding Engineer

Dec 2022 - Jul 2024

UniFind Networks LLP - AI Recruiting Platform

Hyderabad, India

- Owned product and technical delivery from recruiter discovery and first-principles scoping to architecture, Python backend services, REST APIs, document extraction, scoring, ranking, and deployment.
- Scaled candidate-document processing to 1,000+ files and reduced manual shortlisting time by approximately 70%; iterated on customer feedback and production failures.
- Led demonstrations, integrations, release support, and troubleshooting; converted recurring needs into reusable workflow and data-processing capabilities.

Principal Consultant - HCM

Jul 2021 - May 2022

PeopleStrong Technologies Pvt Ltd.

Gurgaon, India

- Supported production HCM and payroll systems for 11+ multinational customers, translating complex requirements into rules, reports, integrations, and workflow automation.
- Developed and debugged 30+ customer-specific rules, reports, integrations, and automations; used SQL in analysis and coordinated incident resolution with product and engineering teams.

PROJECTS

MARS: Multi-Agent Research Intelligence System - Python, FastAPI, LangGraph, Qdrant, React, Docker

- Built and deployed a seven-agent document-research product for planning, retrieval, tool use, synthesis, question answering, citation validation, and reporting; orchestrated state, routing, and recovery with LangGraph.
- Implemented authenticated async APIs, document ingestion, chunking, vector search, grounded context, structured responses, tests, logging, and retrieval/answer evaluation. Live: mars-ai.app

Low-Latency Voice Agent - Python, Streaming APIs, Speech Recognition, LLM Tool Use

- Developing a real-time voice-agent prototype with streaming speech, turn handling, contextual responses, API actions, and end-to-end latency and task-success measurements.

EDUCATION

M.S. Computer Engineering - Machine Learning & Intelligent Systems

Aug 2024 - May 2026

The George Washington University, Washington, DC

Coursework: Advanced Machine Learning, Pattern Recognition & ML, Machine Intelligence, Reinforcement Learning, Big Data & Cloud Computing, Network Security

TECHNICAL SKILLS

Backend & Data: Python, SQL, FastAPI, Async/REST APIs, Pydantic, Data Pipelines, ETL, Document Processing, PostgreSQL, MongoDB, Redis, Pandas, NumPy, Git/GitHub

Agentic AI & Retrieval: LangGraph, LangChain, Multi-Agent Orchestration, Tool/Function Calling, LLM APIs, Prompt & Context Engineering, RAG, Qdrant, FAISS, Embeddings, Semantic Search

Applied ML: PyTorch, TensorFlow, scikit-learn, Transformers, NLP, Model Evaluation, Experiment Design, Data Validation, Retrieval Benchmarks, Agent Regression Tests

Production & Delivery: AWS, Azure, Docker, Kubernetes, CI/CD, Authentication/RBAC, Logging, Monitoring, Latency/Cost Analysis, Failure Recovery, Customer Discovery, Product Scoping, Technical Demos